

Basics of database handling, ADO.NET

Software techniques




Department of
Automation and
Applied Informatics

Agenda

- Fundamentals of relational data models
- SQL
- Object relational mapping
- ADO.NET

Overview, architecture

- The scope of this lecture
 - > Next semester a whole class deals with databases
 - > This is just an overview now
 - > Goal: be able to develop applications that use database.
- Most database management systems support **relational** models
 - > Best known DBMS products: Oracle, IBM DB2, Microsoft SQL Server, MySQL
 - > Well suited for structured data
- Relationless (**NoSQL**) databases are growing in popularity
 - > Most scale very well to large record sizes. They tend to handle a lack of structure in the data better.

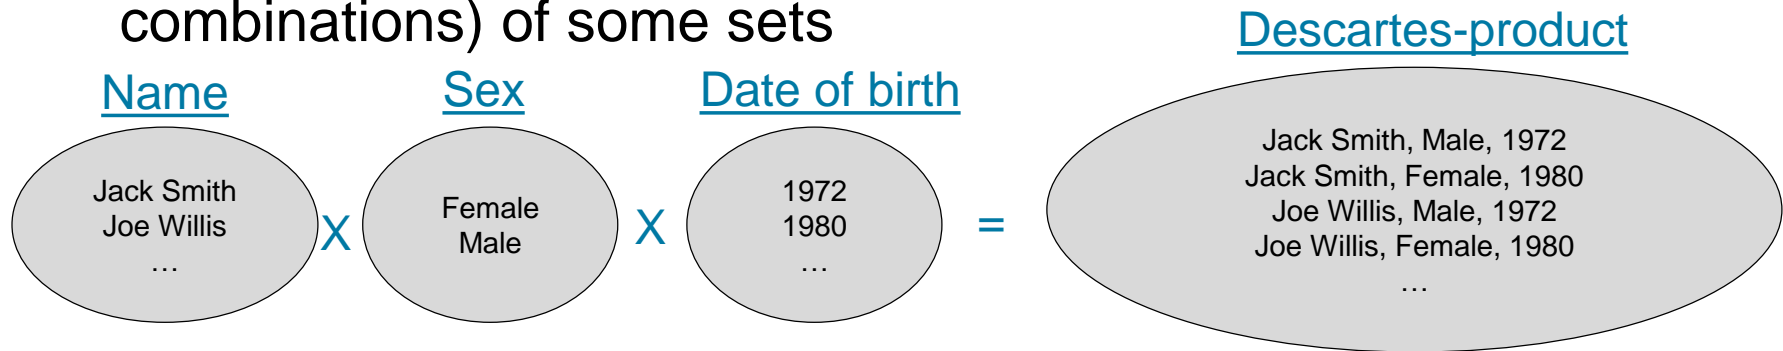


The fundamentals of relational data model

(Short excerpt from the databases course)

Fundamentals of relational data models

- Most database management systems support relational models
- Relation: a subset of the Descartes-product (all the possible combinations) of some sets



- In other words, a relation is a 2-dimensional table with named columns and any number of rows.

<u>Name</u>	<u>Sex</u>	<u>Date of birth</u>
Jack Smith	Male	1972
Joe Willis	Male	1980
...

Basic concepts of relational data models

- Attribute

- > The columns of the table
- > In the previous example there are three attributes: Name, Sex and Date of birth
- > Has got a domain (e.g. bool: true/false, 32 bit int, max 200 char long string, ...)

- Record

- One row of the table, also called an element of the relation

- Relational schema

- Describes the attributes of a relation:
 - RelationName($A_1, A_2, \dots A_N$)
 - Example: Person(Name, Sex, Date of birth)

Concepts of relational data models

- There is a single item in every cell of the table
 - > Each row is unique
 - > Each column has a unique name
 - > The order of rows is not important
 - > The order of the columns is not important
- Relational database schema: the set of relation schemas of each table

Fundamentals of relational data models

■ Anomalies

Example: every student choses a specialization that belongs to a department

<u>Neptun</u>	<u>Name</u>	<u>Year</u>	<u>Specialization</u>	<u>Dep.name</u>	<u>Specialization manager</u>
SD3ER5	George Gee	4	Infocomm	HIT	Dr. Tranzisztor Béla
AD1212	Joe Smith	3	Electrical eng.	VET	Dr. Apertúra János
DHJ7M4	Stephen Willis	3	Software eng.	AUT	Dr. Reláció György
96JGTG	Jane Ewing	4	Software eng.	AUT	Dr. Reláció György
85JUHG	Peter Blacksmith	3	Software eng.	AUT	Dr. Reláció György

There is redundancy in the data

- > **Update anomaly:** if the value of an attribute is changed it has to be modified in more than one places (e.g. modifying the faculty manager for AUT)
- > **Delete anomaly:** if I delete Joe Smith the data about the whole faculty is lost
- > **Insert anomaly:** a new specialization cannot be inserted without inserting a new a student.

Basic concepts

- Functional dependency

- > If the rows of relation R match on attributes A^1, A^2, \dots, A^n , and this also entails that they match on attribute B, then B is functionally dependent on A^1, A^2, \dots, A^n
- > In simpler terms: Attribute B is functionally dependent on A^1, A^2, \dots, A^n , if knowing their value we also know B's value.
- > Notation: $A^1, A^2, \dots, A^n \rightarrow B$

- Examples

- > NeptunCode \rightarrow Name
- > Specialization \rightarrow Department, Specialization manager
- > Name, Address, DoB \rightarrow Social Security Number (SSN)

Basic concepts

- Key

- > A set of attributes that must meet two criteria:

- 1. All other attributes must be functionally dependent on the key: **unique**.
 - For each row, the key must be unique.
 - The key's value identifies the row.
 - 2. If we remove any of the attributes from the set, the first criteria will no longer hold: **minimal**.

- Example:

- > Student(SSN, NeptunCode, Name, Address, DoB)

- Func deps: SSN → NeptunCode, Name, Address, DoB
NeptunCode → SSN, Name, Address, DoB
Name, Address, DoB → SSN, NeptunCode
 - Keys: (SSN), (NeptunCode), (Name, Address, DoB),
because they all determine a single row (they are
unique in the relation)

Basic concepts

- **Primary key (PK)**

- > We select one of the keys for this purpose
- > We underline it in notations: **Student**(NeptunCode, SSN, Name, Address, DoB)

- **Foreign key (FK)**

- > A reference to another table's primary key
- > Example:
 - Specialization**(SpecName, Department, SpecManager)
 - Student**(NeptunCode, Name, **SpecName**, Year)
 - The SpecName attribute in the Student table is a foreign key to elements in the Specialization table.
- > We describe relations between tables with foreign keys
- > Database management systems (DBMS) check the value of foreign keys
 - Eg. A Student may only have a SpecName that appears in the Specialization table
 - As such foreign keys are important to maintaining data integrity!

Basic concepts

- NULL attributes
 - If specified, some attributes can also have NULL value.
- Details about primary keys
 - > In practice we usually generate an 'artificial' primary key
 - > Eg.
 - Employee (EmployeeID, Name, Address, ...)
 - Comodity (ComodityID, Name, Price, Description,...)
 - > Usually one of the two:
 - AutoIncrement ed - the DBMS increases its value automatically by one each time a new row is inserted.
 - GUID (Globaly Unique Identifier) – a 128 bit randomly generated value

Normalization

- Normalization

- > Some table schemas lead to anomalies
- > To avoid this, we decompose the relation (table) to multiple relations, taking the shape of some **normalized form**
- > In practice we use **BCNF** and **3NF**
 - Out of the scope of this lecture
- > Let's normalize our previous example
 - Specialization(SpecializationName, Depart., SpecManager)
 - Student(NeptunCode, Name, Year, **SpecName**)

Student table

NeptunCode	Name	Year	SpecName
SD3ER5	Kotányi Aladár	4	Infokomm. rendszerek
AD1212	Kalapács Attila	3	Vill. energetika
DHJ7M4	Lopovszky Károly	3	Szám. rendszerek
96JGTG	Keltai János	4	Szám. rendszerek
85JUHG	Baldai Balvin	3	

SpecName is a foreign key in the Student table

Specialization table

SpecializationName	Depart.	SpecManager
Infokomm. rendszerek	HIT	Dr. Tranzisztor Béla
Vill. energetika	VET	Dr. Apertúra János
Szám. rendszerek	AAIT	Dr. Reláció György

Normalization

- Denormalization

- > We sometimes (rarely) 'break' a normalized schema
 - Eg. To store a computed value
 - Eg: the number of students could be stored in the Specialization table (to avoid counting each time)
- > Only if we must!
 - Hard to properly maintain, can easily lead to inconsistency!

- Transactions


- > Execute several operations in one go
- > Maintain ACID principles
 - Atomicity, Consistency, Isolation, Durability
 - (Later in other courses)

Object- relational mapping

ORM

Object-relational mapping


- We design databases in two main ways
 - > A) „Traditional”: we plot and decompose relations, normalizing them. This is what we saw before.
 - > B) Take a theoretical entity model (eg. UML class diagram or Entity-Relationship diagram) → Classes/entities will become tables.
- Example for B)
 - > Take the following UML class diagram



- > Looking back at the first lecture – mapping connections in code:
 - **One-to-many connection - Lead professor to Subject:** The LeadProfessor class has a list of pointers or references to the Subject objects they are in charge of. The Subject class has a pointer or reference to its LeadProfessor.
 - **Many-to-many connection - Student – Subject:** both classes have a list of pointers/references to the other.


Object-relational mapping

- How do we get from the class diagram to database tables.

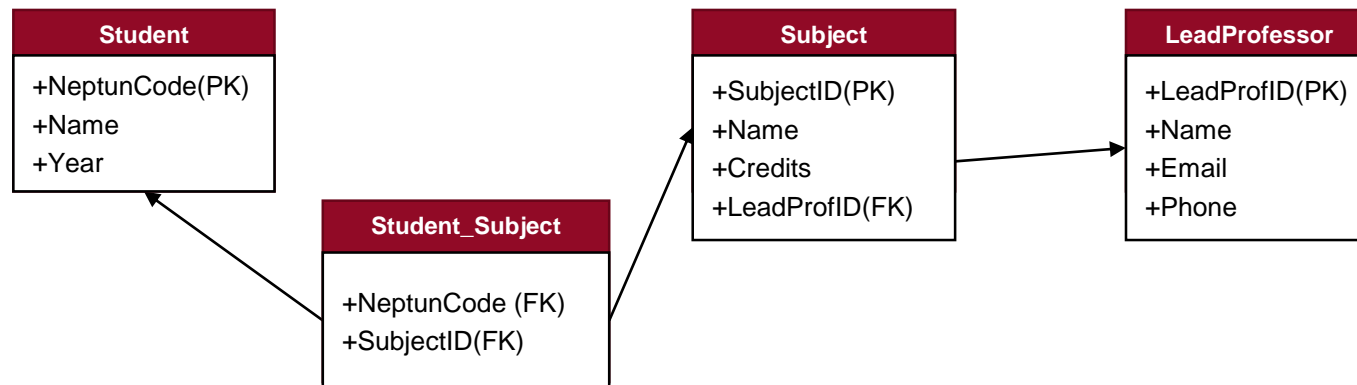


- **Classes, member variables:**
 - Tables, attributes
- **One-many relationship**
 - on the many side of the relation introduce a foreign key into the table
- **Many-many relationship**
 - introduce a new joining table that contains a foreign key for both tables

Object-relational mapping



- This is the database schema:



> We have a fourth **connecting table**: Student_Subject

Object-relational mapping

- Let's look at our database with some data

Student		
<u>NeptunCode</u>	<u>Name</u>	<u>Year</u>
SD3ER5	John Doe	4
AD1212	Kalapács Attila	3
DHJ7M4	Lopovszky Károly	3
96JGTG	Keltai János	4

Subject			
<u>SubjectID</u>	<u>Name</u>	<u>Credits</u>	<u>LeadProfID</u>
1	Prob. theory	2	2
2	Software te.	4	1
3	Databases	3	1

Student_Subject	
<u>NeptunCode</u>	<u>SubjectID</u>
SD3ER5	2
SD3ER5	3
AD1212	1
...	...

John Doe is a student of Szoftech and Database subjects.

Dr. Reláció György is lead professor for Szoftech and Databases.

LeadProfessor		
<u>LeadProfessorID</u>	<u>Name</u>	...
1	Dr. Reláció György	
2	Kalapács Attila	
3	Lopovszky Károly	
4	Keltai János	

Check: are there anomalies?

The basics of SQL


A Structure Query Language

SQL introduction

- **SQL - Structured Query Language**
- Language elements can be split into two groups
 - > **DDL - Data Definition Language**
 - E.g. Creating and modifying tables
 - > **DML – Data Manipulation Language**
 - E.g. Inserting, deleting, selecting, updating records

- **Example**

- > Product and Manufacturer: a product is produced by one manufacturer, a manufacturer produces multiple products
- > Manufacturer(ManufacturerID, Name, Address)
- > Product(ProductID, Name, Price, ManufacturerID)



SQL – data definition

■ Creating the tables

```
CREATE TABLE Manufacturer
(
    ManufacturerID int NOT NULL PRIMARY KEY,
    Name nvarchar(50) NOT NULL,
    Address nvarchar(200) NOT NULL
)

CREATE TABLE Product
(
    ProductID int NOT NULL PRIMARY KEY,
    Name nvarchar(50) NOT NULL,
    Price int NOT NULL,
    ManufacturerID int NULL REFERENCES Manufacturer(ManufacturerID)
)
```

- CREATE TABLE <table name>, and then comes the list of columns. For every column:
 - Name, type, null is allowed or not (optional)
 - optional: PRIMARY KEY (the column is the primary key of the table), REFERENCES <referenced table name>(<referenced column name>) – for defining foreign keys

SQL – data manipulation

Insert a row into a table

```
INSERT INTO Manufacturer(ManufacturerID,Name,Address) VALUES(1,'Nike','Beaverton')
INSERT INTO Product(ProductID, Name, Price, ManufacturerID) VALUES (1,
'Nike Flyknit Air Max', 250, 1)
```

INSERT INTO <table name>(list of columns) VALUES (list of values)

Delete (some) row(s)

```
-- Delete all products
DELETE FROM Product
-- Delete all products that cost 400
DELETE FROM Product WHERE Price = 400
```

DELETE FROM <table name> (column list) WHERE <condition>

Update (some) row(s)

```
-- update product with id = 2
UPDATE Product SET Name = 'Nike Air Max 2014', Price = 220 WHERE ProductID = 1
-- Increase the price of all product that cost less than 300
UPDATE Product SET Price = Price*1.2 WHERE Price < 300
```

**UPDATE <table name> columnname1=value1, columnname2=value2
WHERE <condition>**

SQL queries

■ Query

SELECT list of columns FROM table name WHERE <condition>

■ Conditions

- =, <, >, >=, <=, <>
- exp1 BETWEEN exp2 AND exp3
- exp1 LIKE <string pattern>, the string pattern may contain the following joker characters:

- _: any character
- %: any string

Examples:

- Strings starting with „bob” : LIKE ‘bob%’
- Strings containing „bob” : LIKE ‘%bob%’
- exp IS NULL
- exp IS NOT NULL
- AND, OR, NOT
- EXISTS, NOT EXISTS: a result set contains any record or not

SQL queries

■ Examples

```
-- Query every record of  
-- the product table  
SELECT *  
FROM Product
```



	TemekID	Nev	Ar	GyartoID
1	1	Termék1	100	1
2	2	Termék2	200	1
3	3	Termék3	400	2
4	4	Termék4	400	NULL

```
-- Query every product  
SELECT Name, Price  
FROM Product
```

```
SELECT *  
FROM Product  
WHERE Price > 400
```

```
SELECT *  
FROM Product  
WHERE Price > 400 OR Price < 200
```

```
SELECT *  
FROM Product  
WHERE Price BETWEEN 200 AND 300
```

```
SELECT *  
FROM Product  
WHERE ManufacturerID IS NULL
```

```
SELECT *  
FROM Product  
WHERE Name LIKE '%unar%'
```

SQL queries

■ Example 1: List all manufacturers with all their products

```
SELECT Manufacturer.*, Product.Name AS ProductName
FROM Manufacturer JOIN Product ON
    Manufacturer.ManufacturerID = Product.ManufacturerID
```

	GyartoID	Nev	Cim	TemekNev
1	1	Gyártó1	Cím1	Termék1
2	1	Gyártó1	Cím1	Termék2
3	2	Gyártó2	Cím2	Termék3

- > JOIN, or INNER JOIN: joins two tables using the condition after the ON keyword
 - For every record in the manufacturer table it searches for the products where `Manufacturer.ManufacturerID = Product.ManufacturerID` is true.

SQL queries

- Example 2: List all the products with the corresponding manufacturer name

```
SELECT Product.*, Manufacturer.Name AS ManufacturerName
FROM Product JOIN Manufacturer
    ON Product.ManufacturerID = Manufacturer.ManufacturerID
```

- Example 3: List all the products with the corresponding manufacturer name, but also list the products where manufacturer is not specified (ManufacturerId is null)

```
SELECT Product.*, Manufacturer.Name AS ManufacturerName
FROM Manufacturer RIGHT OUTER JOIN Product
ON Product.ManufacturerID = Manufacturer.ManufacturerID
```

- A LEFT OUTER JOIN takes the left side record into the result set even if there is no matching right side record
- A RIGHT OUTER JOIN takes the right side record into the result set even if there is no matching left side record

SQL queries

- List all manufacturers that have got at least two products with a price greater than 200

```
select p.ManufacturerId, count(*) as prodCount
from Manufacturer m
RIGHT OUTER JOIN Product p ON p.ManufacturerID = m.ManufacturerID
where p.price > 200
group by p.ManufacturerId
having count(*) > 1
```

Group by

- Groups are created so that in each group the group by parameters have got the same value
- The select list can only contain columns that are either in the group by clause or in an aggregate function
- Having is used to filter for groups

ADO.NET

ADO.NET

- What is ADO.NET?
 - > Relational database access from .NET environment
 - E.g. Oracle, MSSQL, Access, ...
- What can it be used for?
 - > Execute SQL commands and read the result set
- Other DB access solutions in .NET
 - > Entity Framework is an ORM, much newer technology
 - > ADO.NET is the 'traditional' way of accessing a database, we'll be looking at this
- Approach
 - > The approach of ADO.NET is quite similar to many non-.NET solutions.

ADO.NET – object model

- Most important objects for data access

- > Connection:

- Represents a database connection to a database manager
 - Before executing any command a connection has to be created.

- > Command:

- Can execute a SQL command (or stored procedure)

- > DataReader

- Can be used to iterate through the result set of a query (e.g. SELECT).

ADO.NET – Provider-based architecture

- Provider-independent interfaces and abstract base classes
 - > `IConnection`, `DbConnection`
 - > `ICommand`, `DbCommand`
 - > `IDataReader`, `DbDataReader`
- Provider-specific implementations
 - > Derive from the above mentioned base classes
 - > E.g...:
 - Oracle: `OracleConnection`, `OracleCommand`, `OracleDataReader`
 - MSSQL: `SqlConnection`, `SqlCommand`, `SqlDataReader`
 - ...

ADO.NET – Connection oriented data access

- Steps(using MS SQL provider)
 - > Instantiate `SqlConnection` instance with the parameters of the connection (connection string)
 - > Create an `SqlCommand` instance
 - Specify the SQL command text or stored procedure name as parameter
 - Set the `SqlConnection` instance to use when executing the command
 - > Execute the command (`SqlCommand`)
 - > If the command has got a result set (e.g. SELECT), an instance of `SqlDataReader` is used to iterate through the result set.
 - > Close the connection
- Methods of a command
 - > `ExecuteNonQuery` – for SQL commands that doesn't return anything (e.g. INSERT, UPDATE)
 - > `ExecuteScalar` – for SQL commands that return a scalar value.
 - > `ExecuteReader` – for SQL commands that return one or more result sets.

ADO.NET – Connection oriented data access

```
try
{
    // Create connection object
    conn = new SqlConnection("Data Source=(local); Initial Catalog=SzoftechDB; Integrated security =
true");

    // Create database command
    SqlCommand command = new SqlCommand("SELECT ProductID, Name, Price FROM Product", conn);

    // Open connection
    conn.Open();

    // Query and print results
    reader = command.ExecuteReader();
    while (reader.Read())
    {
        Console.WriteLine("{0}\t{1}\t{2}", reader["ProductID"].ToString(),
            reader["Name"].ToString(), reader["Price"].ToString());
    }
}
finally
{
    if (reader != null) reader.Close();
    if (conn != null) conn.Close(); // We must make sure we close the connection!!!
}
```

Example for a query

ADO.NET – Connection oriented data access

- Key features

- > **DataReader** can read the result set only once, forward, record by record (e.g. cannot step back to a previous record).
- > `reader["column name"]` returns the value of a column of the current record
- > The database connection is a heavy resource!!!
 - Close it as soon as possible!
 - For non query operations right after executing the command
 - For queries after reading the result set with DataReader.
 - Why is it an expensive resource?
 - Usually the number of connections that can be opened in parallel is limited: if many clients connect to a DB management system it can run out of free connections.
 - The max concurrent connection number is usually also limited by product licenses. You have to pay for more connections

ADO.NET – Connection oriented data access

```
SqlConnection conn = null;
try
{
    // Connect to the database
    conn = new SqlConnection("Data Source=.\SQLEXPRESS; Initial Catalog = SzoftechDB; Integrated Security = True");

    conn.Open(); // Open the connection

    //Watch out for SQL-INJECTION by passing parameters appropriately!!!
    SqlParameter pProductId = new SqlParameter("@ProductId", productId);
    SqlParameter pName = new SqlParameter("@Name", name);
    SqlParameter pPrice = new SqlParameter("@Price", price);

    // Create the command
    SqlCommand command = new SqlCommand(
        "INSERT INTO Product(ProductId, Name, Price) VALUES(@ProductId, @Name, @Price)");

    command.Parameters.Add(pProductId );
    command.Parameters.Add(pName);
    command.Parameters.Add(pPrice);

    // Provide the connection
    command.Connection = conn;
    int res = command.ExecuteNonQuery();
}
finally { if (conn != null) conn.Close(); }
```

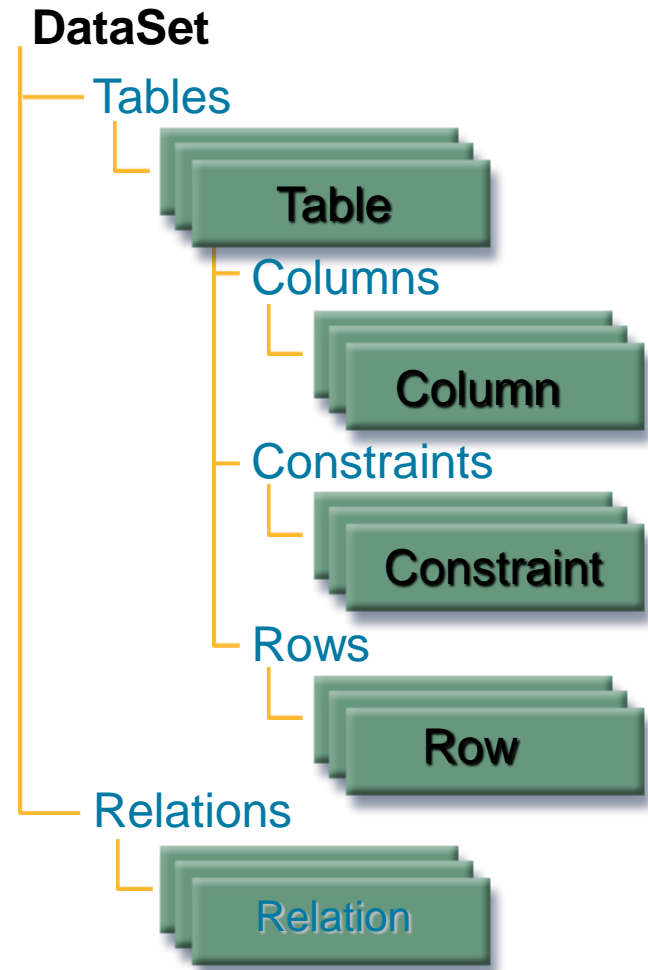
Example for insert

ADO.NET – ‘Connectionless’ data access

- ‘Connectionless’ data access
 - > Aka „DataSet-based” data access
 - > Usual scenario
 - Query data (SELECT) from database and close the connection
 - Display data for the user (e.g. in a datagrid)
 - User edits data
 - User synchronizes the modifications: connects to the database and performs the modifications (INSERT, UPDATE, DELETE)
 - > Problems to solve
 - After querying the data it has to be stored in memory
 - User modifications have to be stored (new rows, edited rows, deleted rows).
 - > Solutions for these problems
 - Custom solutions
 - Using DataSet

ADO.NET – DataSet-based data access

- What is a DataSet?
 - > A temporary in-memory database (tables, columns, rows, relations)
- Key features
 - > The table stores the retrieved data
 - > Keeps track of all the modifications



ADO.NET – DataSet-based data access

- Steps

- 1) Query

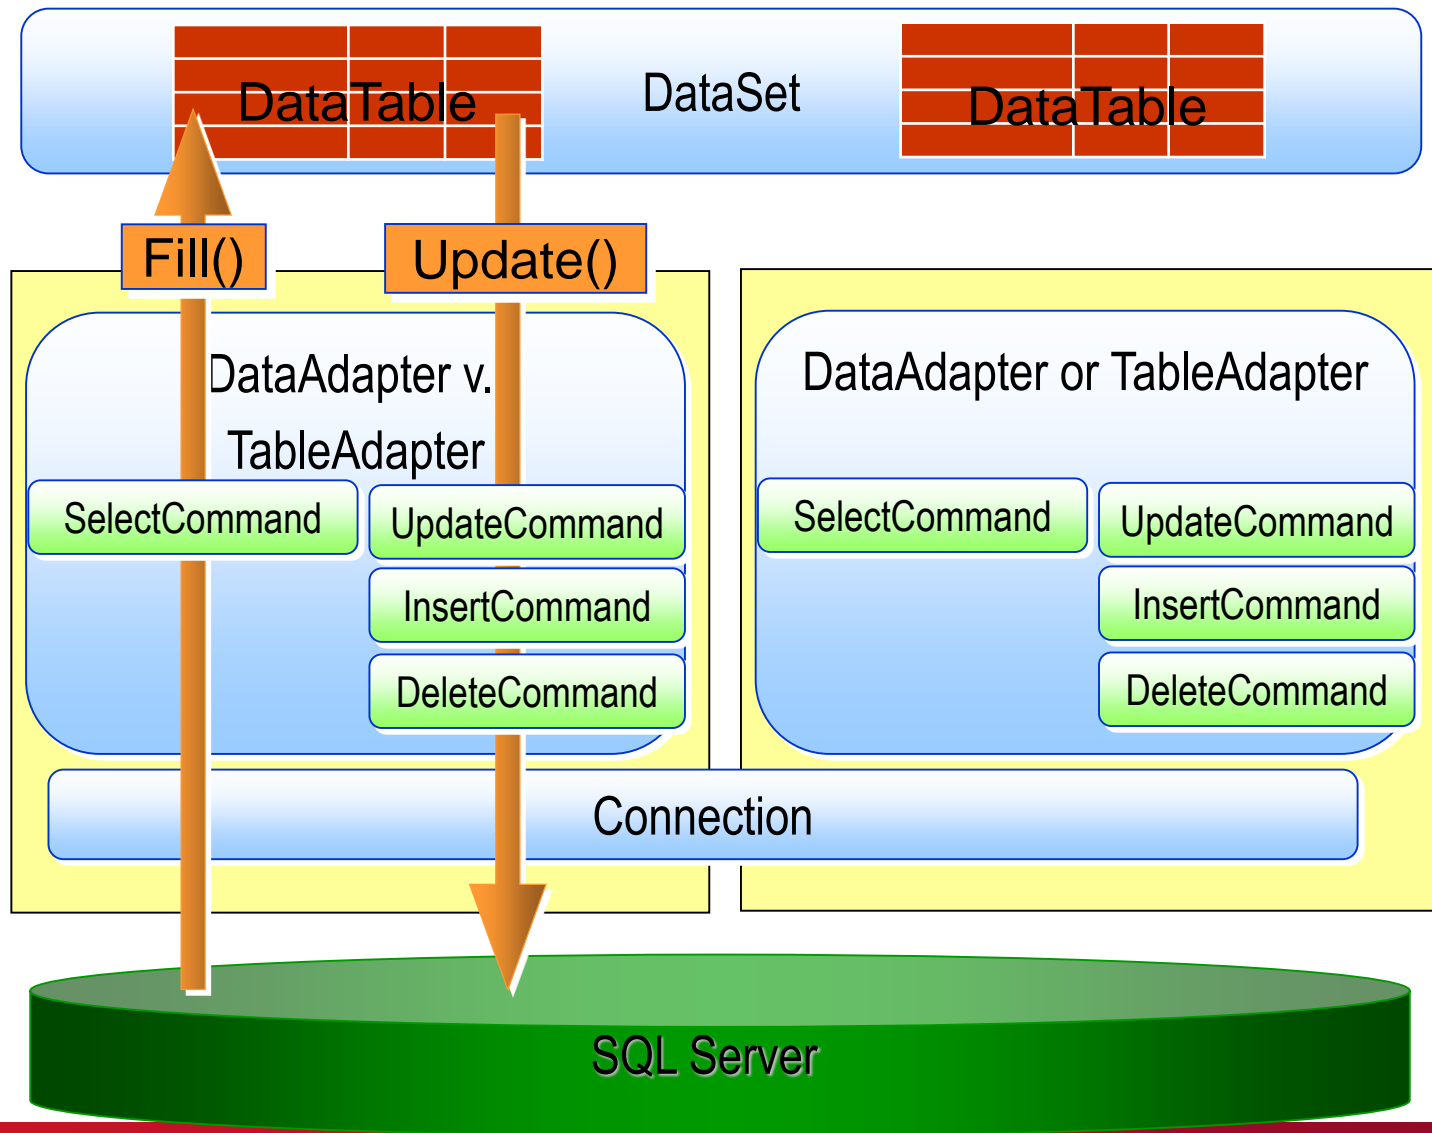
- Open the connection
- Fill the tables of the DataSet from the result of a database query
- Close the connection

- 2) Display data from the DataSet & modify data (in memory)

- 3) Sync modifications back to database

- Open the connection
- Save the modified rows of the DataSet
 - INSERT for every new row
 - UPDATE for every modified row
 - DELETE for every deleted row
- Close the connection

ADO.NET – DataSet based data access



ADO.NET – DataSet based data access

```
// Create connection object
SqlConnection conn = new SqlConnection(
    "Data Source=(local); Initial Catalog=SzoftechDB;...");

// Create DataSet object
DataSet dsSzoftech = new DataSet();

// Create Data adapter: this will fill the DataSet, and save changes to the database.
SqlDataAdapter adapter = new SqlDataAdapter("SELECT * FROM Product", conn);
adapter.InsertCommand = new SqlCommand(< an INSERT command >);
adapter.UpdateCommand = new SqlCommand(< an UPDATE command >);
adapter.DeleteCommand = new SqlCommand(< an DELETE command >);

// Filling the DataSet "Product" table: runs the adapter's SELECT command
// The connection is opened and closed automatically
adapter.Fill(dsSzoftech, "Product");

// Print the Name of the first record of Product table
Console.WriteLine(dsSzoftech.Tables["Product"].Rows[0]["Name"].ToString());

// Set the first record's price to 999
dsSzoftech.Tables["Product"].Rows[0]["Price"] = 999;

// Save changes to the database: the changed row is updated using the Update command
// The connection is opened and closed automatically
adapter.Update(dsSzoftech, "Product");
```

Looking beyond ADO.NET

- Problems
 - > A DataSet still deals with 'raw' SQL and relations
 - > Wouldn't it be nice to write queries in C#, using C# classes...
- ORM: **Entity Framework**
 - > The most popular high level ORM for .NET (Microsoft)
- Query language: **LINQ**
 - > Also works on collections and other iterable types